


# Decision Tree


Mahdi Roozbahani  
Georgia Tech

Adding new feature to the  
existing code 🤔🤔





# Visual Introduction to Decision Tree



Building a tree to distinguish homes in New York from homes in San Francisco

# Decision Tree: Example (2)

	<b>O</b>	<b>T</b>	<b>H</b>	<b>W</b>	<b>Play?</b>
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

Outlook:

Sunny,  
Overcast,  
Rainy

Temperature: Hot,

Medium,  
Cool

Humidity:

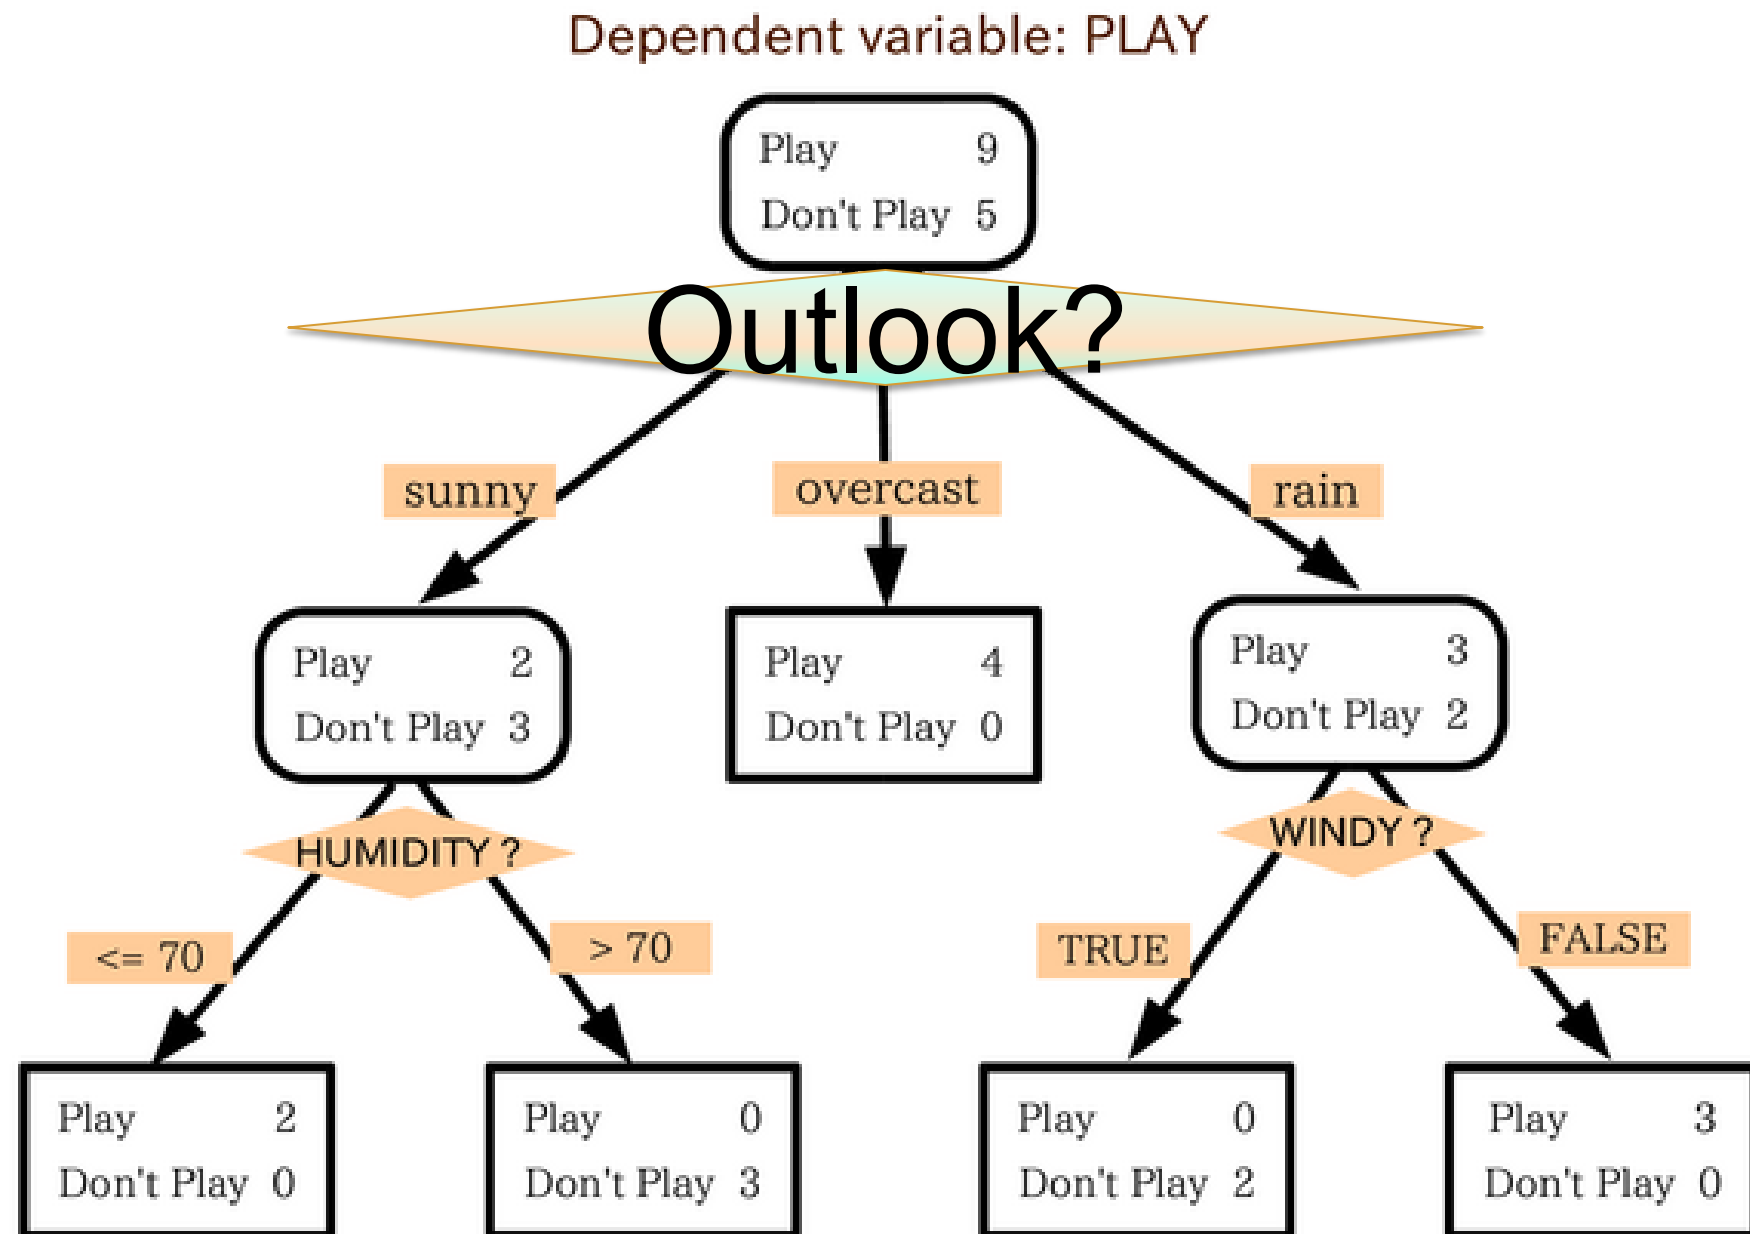
High,  
Normal,  
Low

Wind:

Strong,  
Wweak

*Will I play tennis today?*

# Decision trees (DT)



**The classifier:**

$f_T(x)$ : majority class in the leaf in the tree  $T$  containing  $x$

**Model parameters:** The tree structure and size

# Decision trees

*Attribute = feature = dimension = variable*

Pieces:

1. Find the **best attribute** to split on
2. Find the **best split** on the **chosen attribute**
3. Decide on when to stop splitting


# Categorical or Discrete attributes

- Three variables:
  - Hair = {blond, dark}
  - Height = {tall, short}

Label – Country = {Gromland, Polvia}

Training data:


(B,T,P)  
(B,T,P)  
(B,S,G)  
(D,S,G)  
(D,T,G)  
(B,S,G)





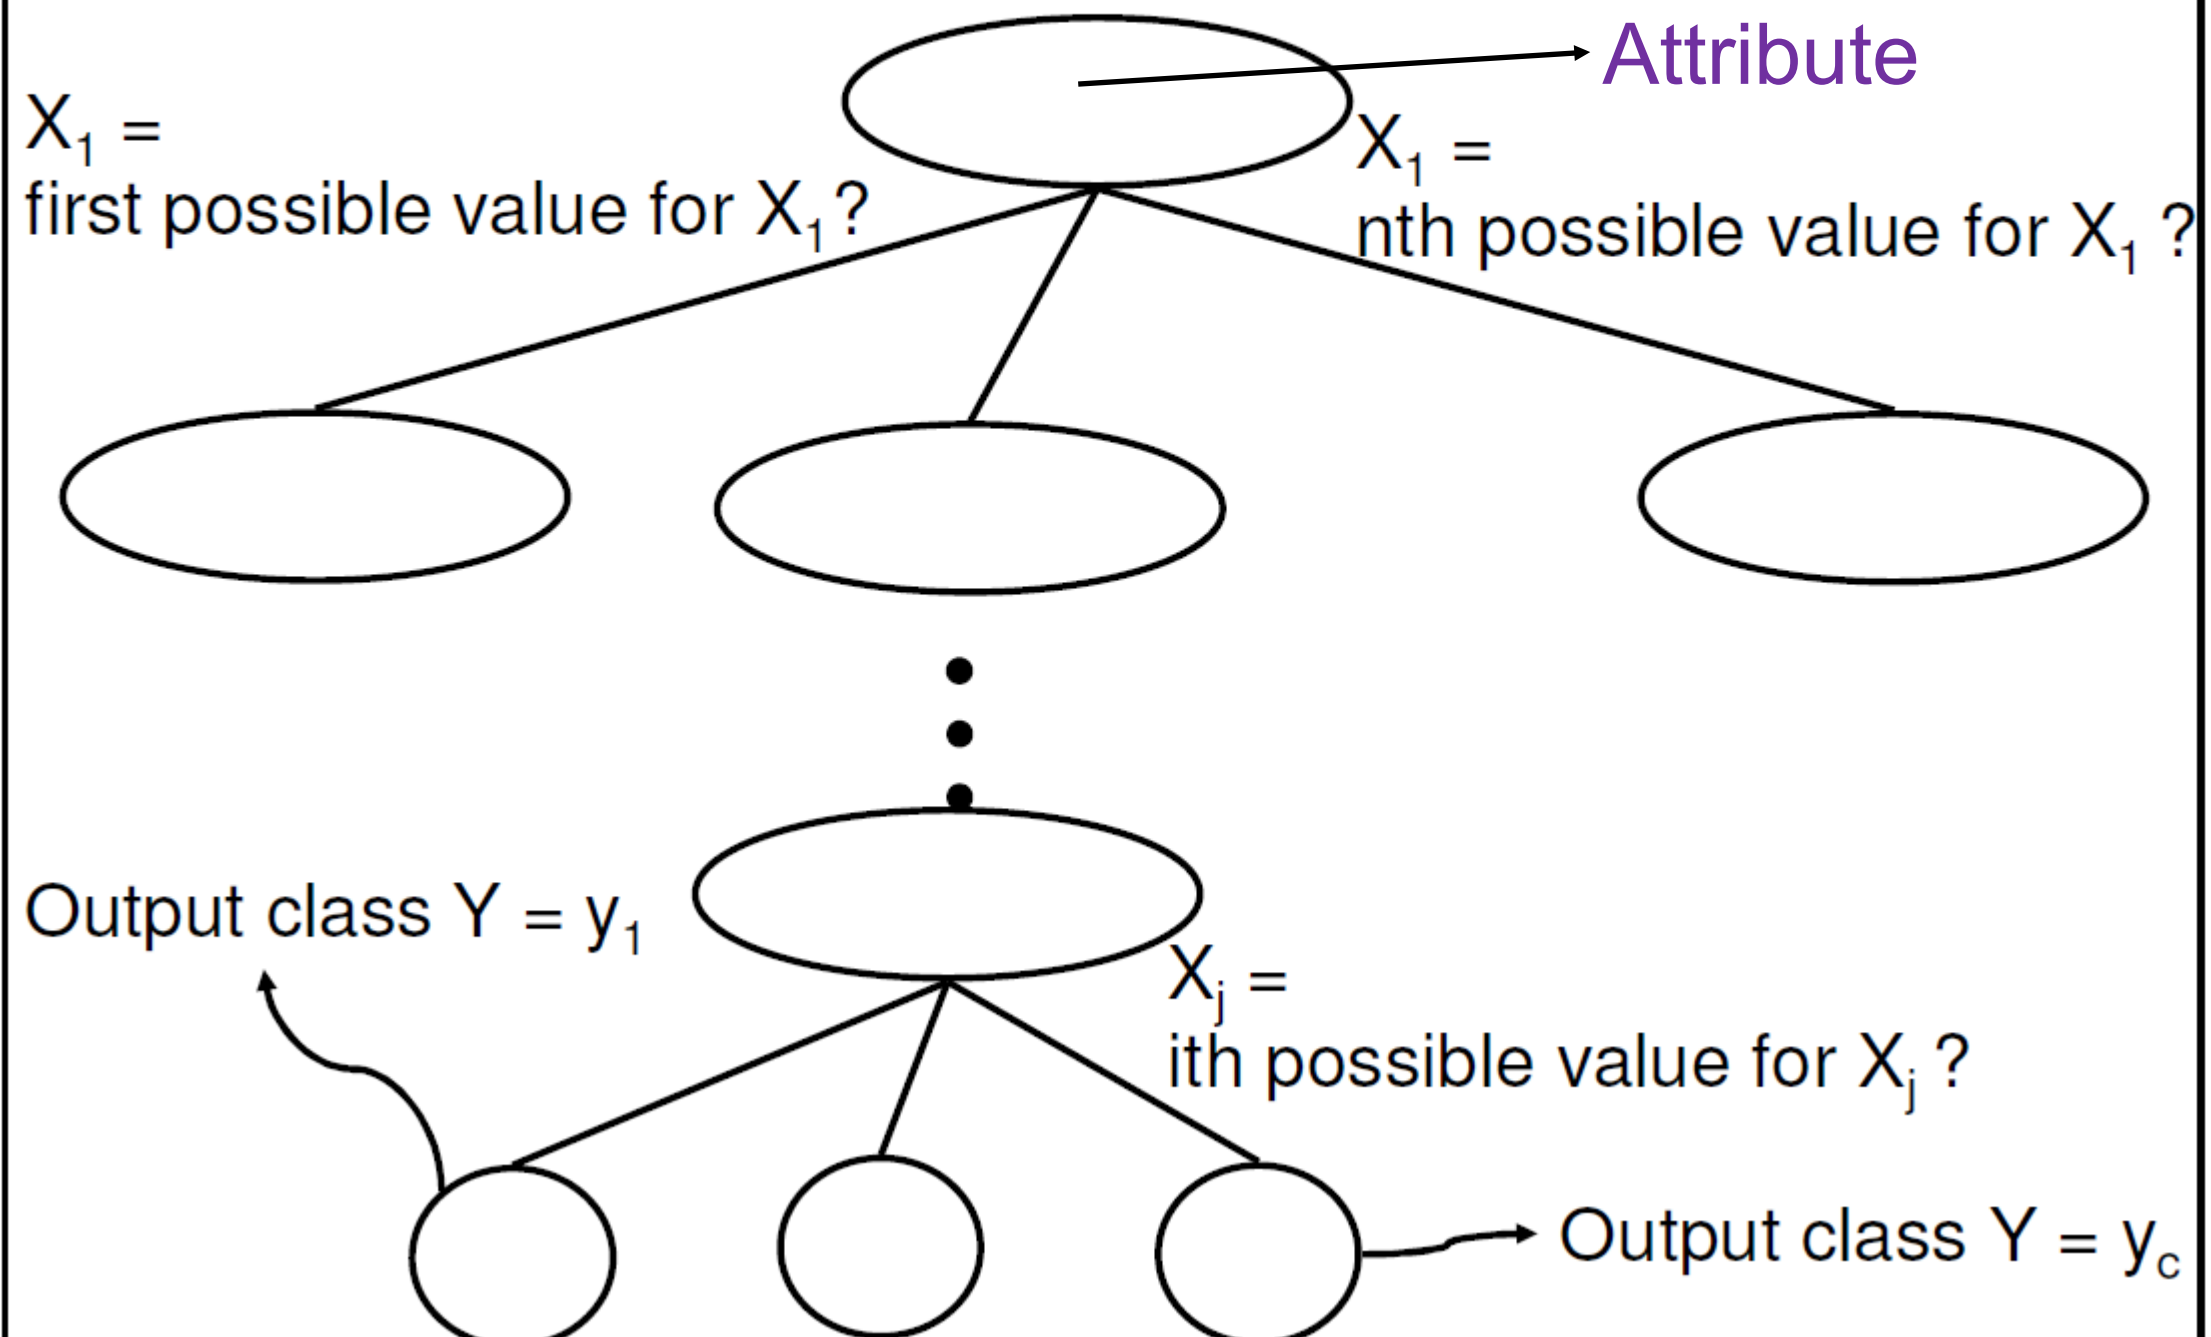
At each level of the tree, we split the data according to the value of one of the attributes

After enough splits, only one class is represented in the node → This is a terminal leaf of the tree  
We call that class the output class for that node



'G' is the output for this node


# General Decision Tree (Discrete Attributes)



# Continuous attributes

## Decision Tree Example





$X_2 = 0.5$



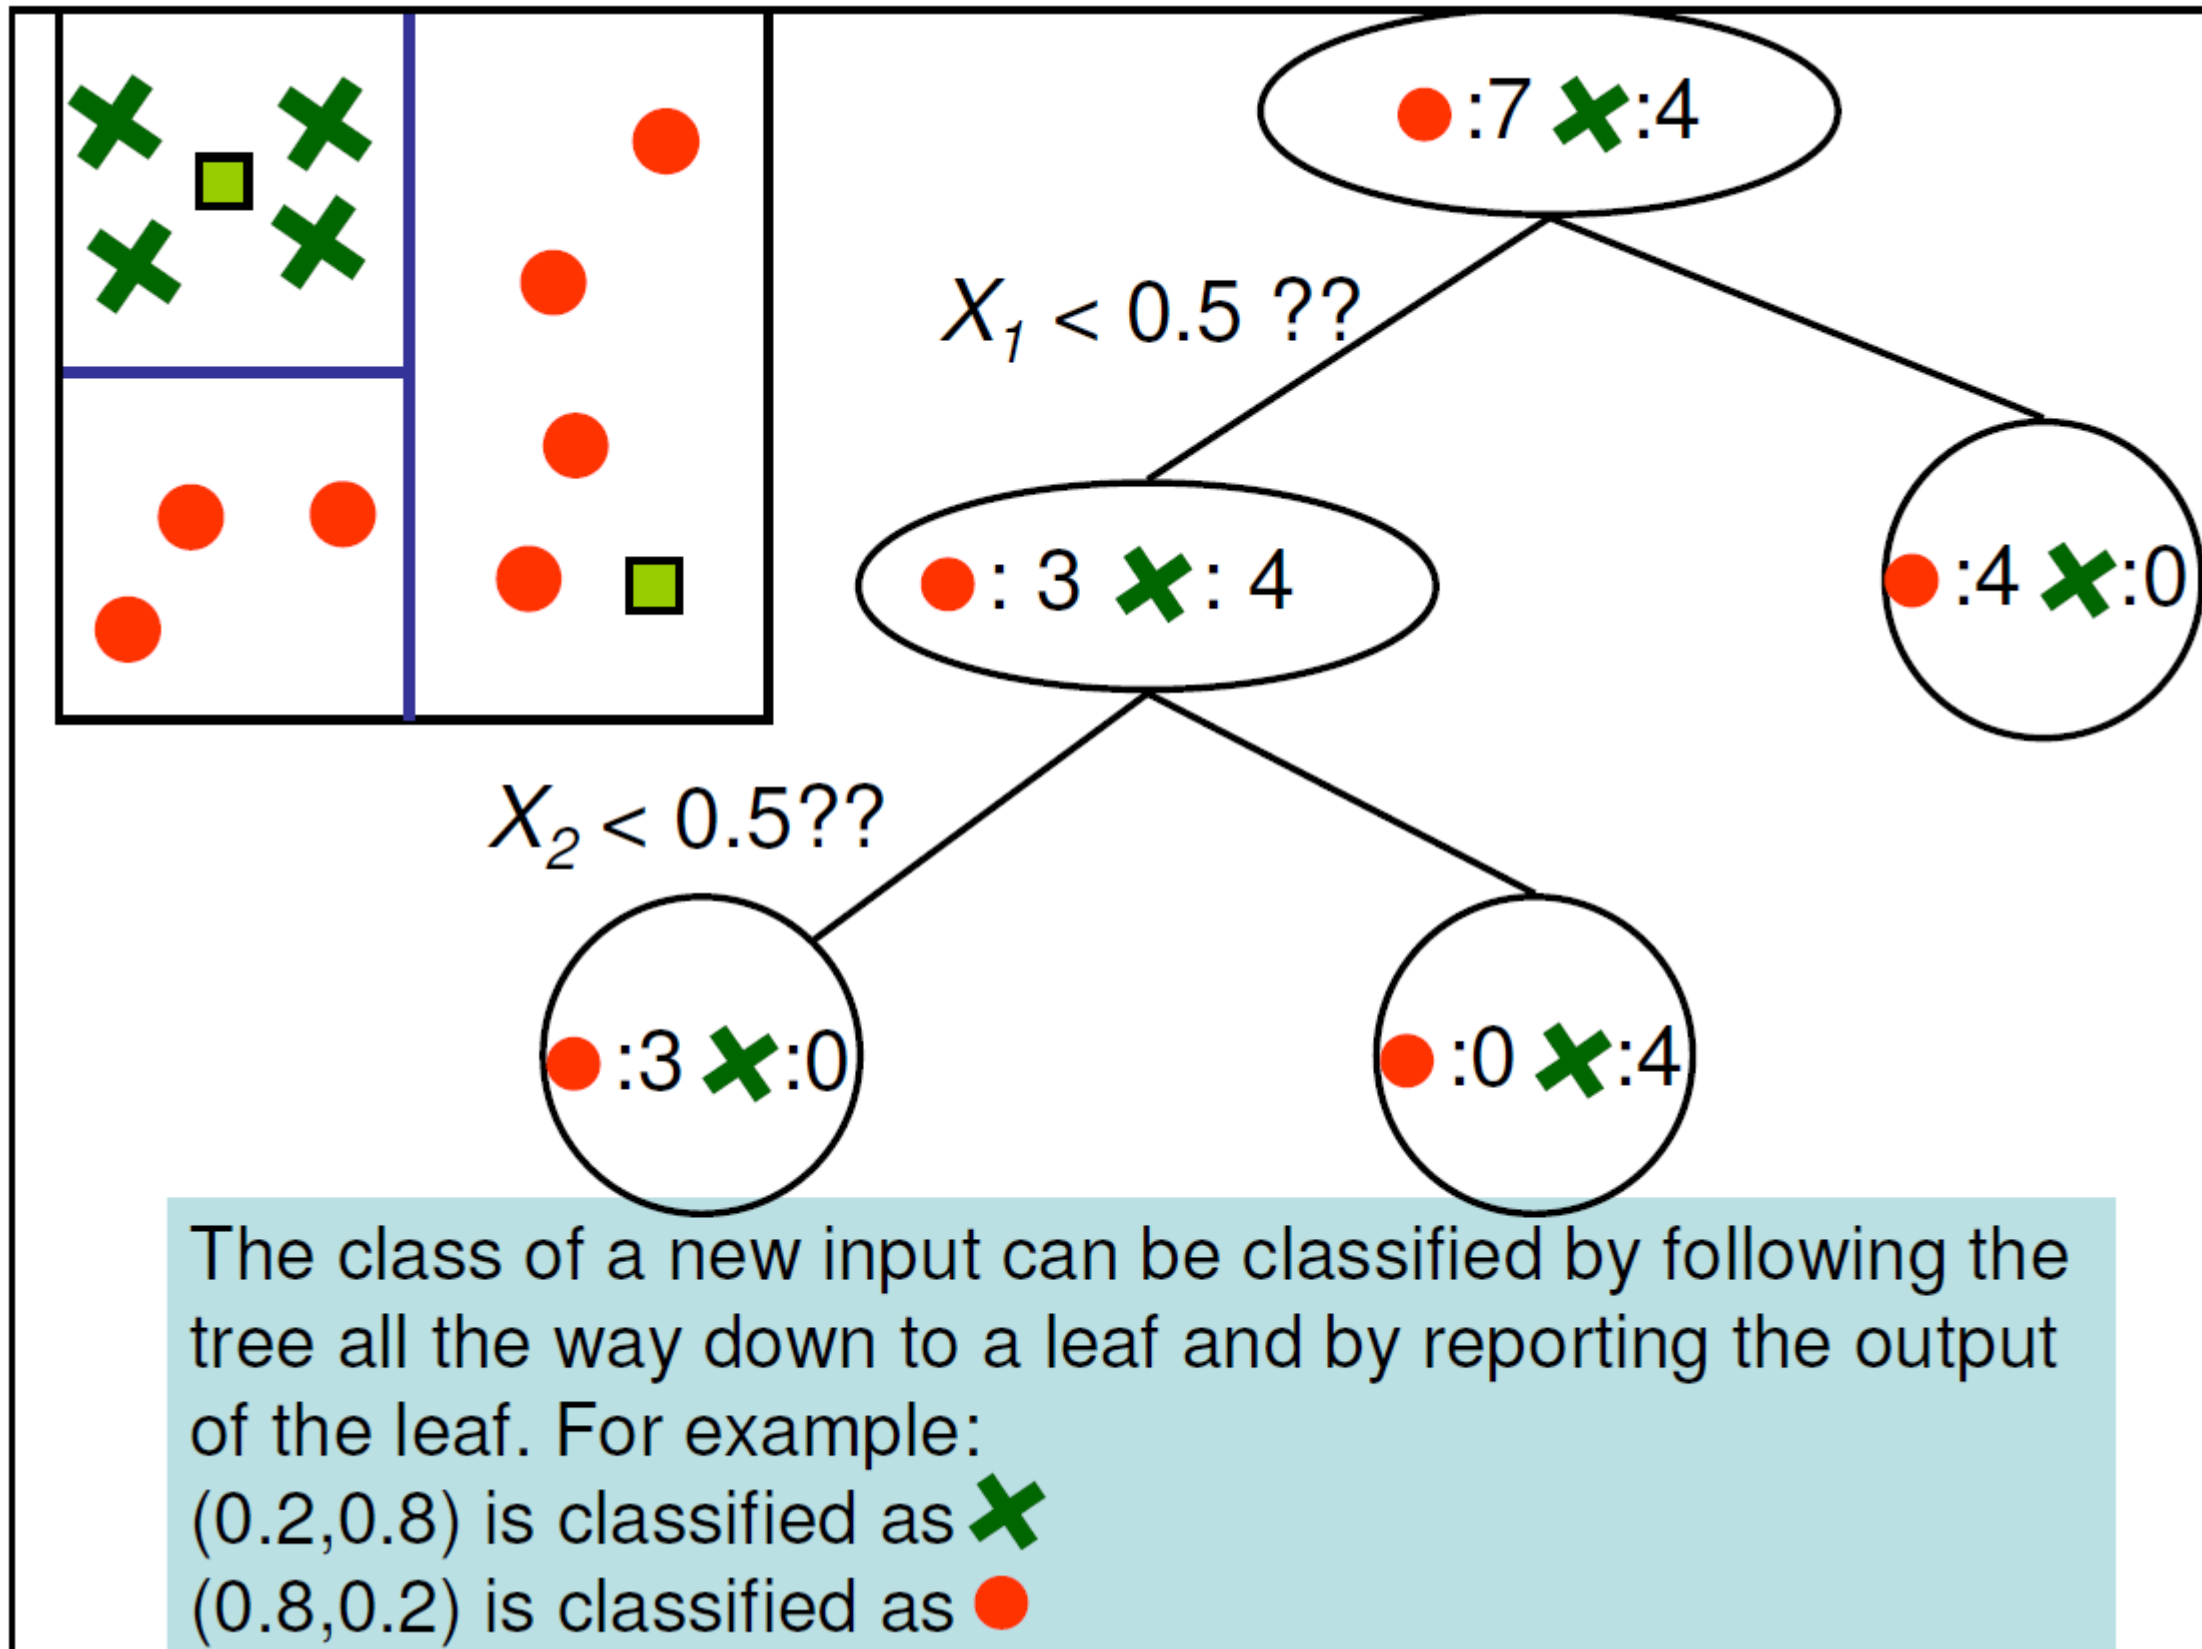
$X_1 = 0.5$

$X_2 < 0.5??$


$X_1 < 0.5 ??$



# Test data




# General Decision Tree (Continuous Attributes)



# Basic Questions

- How to choose the attribute/value to split on at each level of the tree?
- When to stop splitting? When should a node be declared a leaf?
- If a leaf node is impure, how should the class label be assigned?
- If the tree is too large, how can it be pruned?


How to choose the attribute/value to split on at each level of the tree?



- Two classes (red circles/green crosses)
- Two attributes:  $X_1$  and  $X_2$
- 11 points in training data
- Idea  $\rightarrow$  Construct a decision tree such that the leaf nodes predict correctly the class for all the training examples



How to choose the attribute/value to split on at each level of the tree?




Good


Bad





We want to find the most compact, smallest size tree (Occam's razor), that classifies the training data correctly → We want to find the split choices that will get us the fastest to pure nodes



This node is "pure" because there is only one class left → No ambiguity in the class label



This node is almost "pure" → Little ambiguity in the class label



These nodes contain a mixture of classes → Do not disambiguate between the classes

$$\sum p(x) f(x)$$

$$I(x) = \log_2 \frac{1}{p(x)}$$

# Information Content

$$\begin{aligned} H(x) &= \sum p(x) I(x) \\ &= \sum p(x) \log_2 \frac{1}{p(x)} = \\ &= - \sum p(x) \log_2 p(x) \end{aligned}$$

## Coin flip

$C_{1H}$	<b>0</b>
$C_{1T}$	<b>6</b>

$$P(C_{1H}) = 0/6 = 0$$

$$P(C_{1T}) = 6/6 = 1$$

$C_{2H}$	<b>1</b>
$C_{2T}$	<b>5</b>

$$P(C_{2H}) = 1/6$$

$$P(C_{2T}) = 5/6$$

$C_{3H}$	<b>2</b>
$C_{3T}$	<b>4</b>

$$P(C_{3H}) = 2/6$$

$$P(C_{3T}) = 4/6$$

Which coin will give us the purest information? Entropy ~ Uncertainty

Lower uncertainty, higher information gain

$$H(X) = - \sum_{i=1}^N P(x=i) \log_2 P(x=i)$$

$$\text{Entropy} = - 0 \log 0 - 1 \log 1 = - 0 - 0 = 0$$

$$\text{Entropy} = - (1/6) \log_2 (1/6) - (5/6) \log_2 (5/6) = 0.65$$

$$\text{Entropy} = - (2/6) \log_2 (2/6) - (4/6) \log_2 (4/6) = 0.92$$


# Entropy

- In general, the average number of bits necessary to encode  $n$  values is the entropy:


$$H = - \sum_{i=1}^n P_i \log_2 P_i$$

- $P_i$  = probability of occurrence of value  $i$ 
  - High entropy  $\rightarrow$  All the classes are (nearly) equally likely
  - Low entropy  $\rightarrow$  A few classes are likely; most of the classes are rarely observed

# Entropy




The entropy captures the degree of “purity” of the distribution



# Example Entropy Calculation

①




$$N_A = 1$$
$$N_B = 6$$

$$p_A = N_A / (N_A + N_B) = 1/7$$
$$p_B = N_B / (N_A + N_B) = 6/7$$

$$H_1 = -p_A \log_2 p_A - p_B \log_2 p_B$$
$$= 0.59$$

②



$$N_A = 3$$
$$N_B = 2$$

$$p_A = N_A / (N_A + N_B) = 3/5$$
$$p_B = N_B / (N_A + N_B) = 2/5$$


$$H_2 = -p_A \log_2 p_A - p_B \log_2 p_B$$
$$= 0.97$$

$$H_1 < H_2 \Rightarrow (2) \text{ less pure than } (1)$$

$$\sum p(x) f(x)$$

# Conditional Entropy

Entropy before splitting:  $H$



After splitting, a fraction  $P_L$  of the data goes to the left node, which has entropy  $H_L$

After splitting, a fraction  $P_R$  of the data goes to the right node, which has entropy  $H_R$

The average entropy after splitting is:

Entropy of left node


$$H_L \times P_L + H_R \times P_R$$

“Conditional Entropy”

Probability that a random input is directed to the left node



# Information Gain



We want nodes as pure as possible

→ We want to reduce the entropy as much as possible

→ We want to maximize the difference between the entropy of the parent node and the expected entropy of the children

Information Gain (IG) = Amount by which the ambiguity is decreased by splitting the node

Maximize:

$$IG = H - (H_L \times P_L + H_R \times P_R)$$


# Notations

- Entropy:  $H(Y)$  = Entropy of the distribution of classes at a node
- Conditional Entropy:
  - *Discrete*:  $H(Y|X_j)$  = Entropy after splitting with respect to variable  $j$
  - *Continuous*:  $H(Y|X_j, t)$  = Entropy after splitting with respect to variable  $j$  with threshold  $t$
- Information gain:
  - *Discrete*:  $IG(Y|X_j) = H(Y) - H(Y|X_j)$  = Entropy after splitting with respect to variable  $j$
  - *Continuous*:  $IG(Y|X_j, t) = H(Y) - H(Y|X_j, t)$  = Entropy after splitting with respect to variable  $j$  with threshold  $t$




$$p_{\text{red}} = \frac{5}{11} \quad p_{\text{green}} = \frac{6}{11} \quad H = -\frac{5}{11} \log_2 \frac{5}{11} - \frac{6}{11} \log_2 \frac{6}{11} = 0.99$$

$$H = 0.99$$




$$IG =$$

$$H - (H_L * 4/11 + H_R * 7/11)$$




$$H_L = 0$$




$$H_R = 0.58$$

$$H = 0.99$$




$$IG =$$

$$H - (H_L * 5/11 + H_R * 6/11)$$




$$H_L = 0.97$$




$$H_R = 0.92$$


$H = 0.99$



$IG = 0.62$




$H_L = 0$




$H_R = 0.58$


$H = 0.99$



$IG = 0.052$




$H_L = 0.97$




$H_R = 0.92$


$$H = 0.99$$



$$IG = 0.62$$




$$H_L = 0$$




$$H_R = 0.58$$


$$H = 0.99$$



$$IG = 0.052$$




$$H_L = 0.97$$



$$H_R = 0.92$$

Choose this split because the information gain is greater than with the other split


# More Complete Example





● = 20 training examples from class A

× = 20 training examples from class B

Attributes =  $X_1$  and  $X_2$  coordinates



Best split value (max Information Gain) for  $X_1$  attribute: 0.24 with  $IG = 0.138$



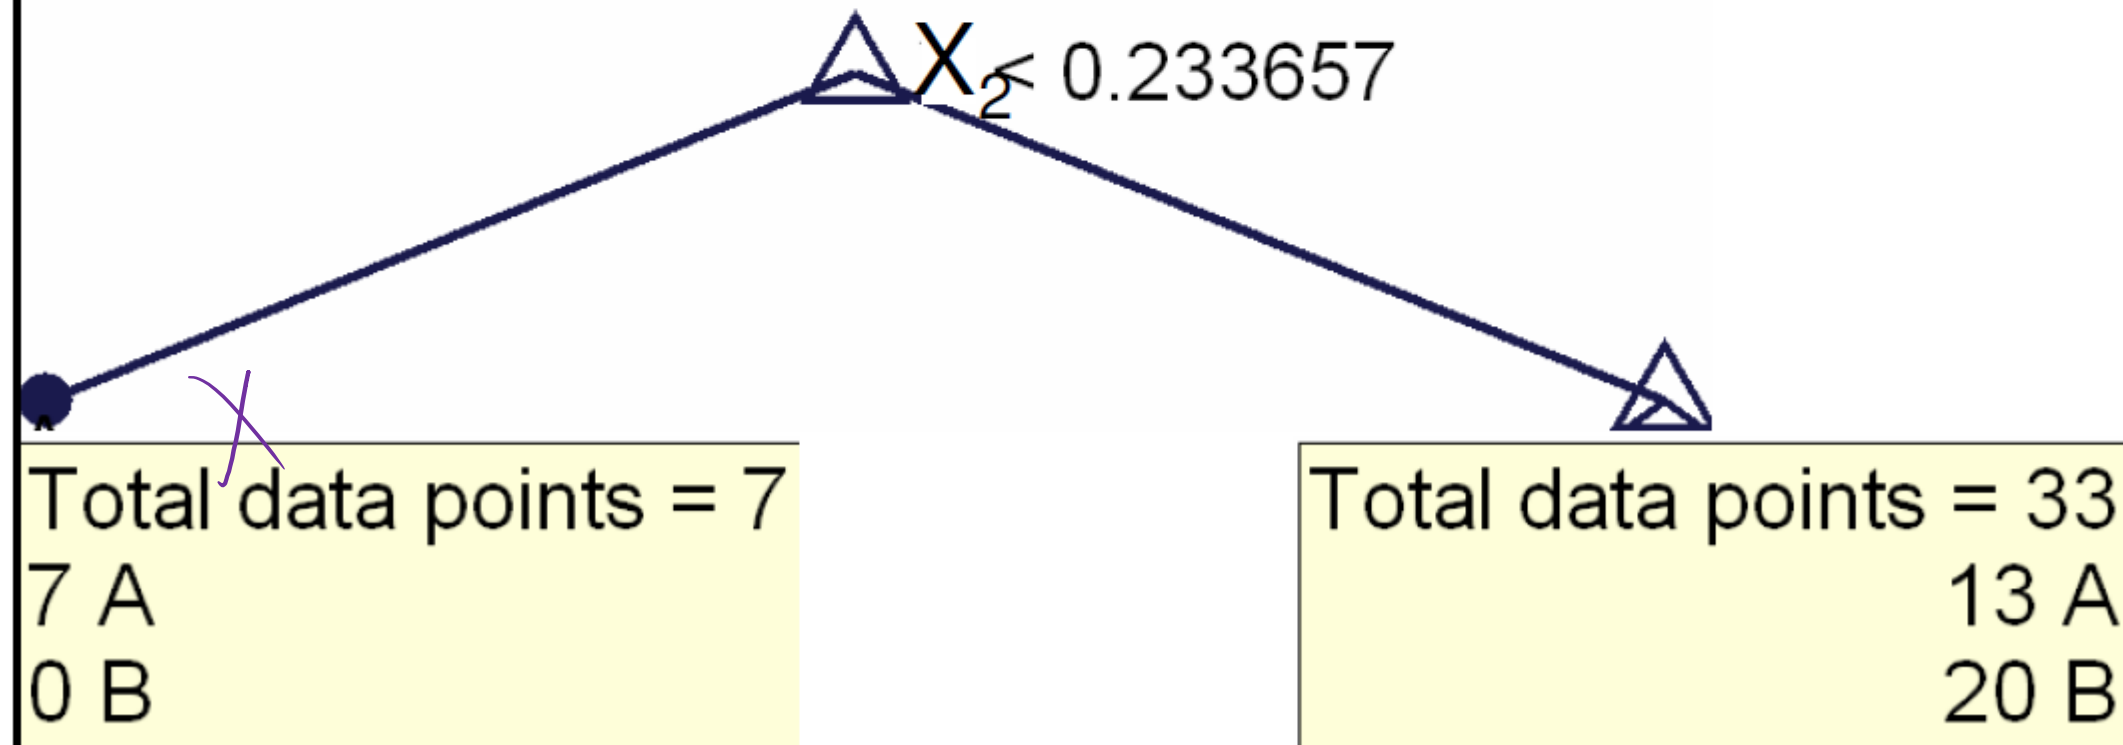

Best split value (max Information Gain) for  $X_2$  attribute: 0.234 with  $IG = 0.202$

Best  $X_1$  split: 0.24, IG = 0.138

Best  $X_2$  split: 0.234, IG = 0.202



Split on  $X_2$  with 0.234




**Left** direction for **smaller** value, **right** direction for **bigger** value

Best X split: 0.24 IG = 0.138

There is no point in splitting this node further since it contains only data from a single class → return it as a leaf node with output 'A'

202



This node is not pure so we need to split further

$X_2 \leq 0.233657$

Total data points = 7

7 A



0 B

Total data points = 33

13 A



20 B





Best split value (max Information Gain) for  $X_1$  attribute: 0.22 with IG  $\sim$  0.182

IG



$X_2$  Split value



Best split value (max Information Gain) for  $X_2$   
attribute: 0.75 with IG  $\sim 0.353$

Best  $X_1$  split: 0.22, IG = 0.182

Best  $X_2$  split: 0.75, IG = 0.353






Split on  $X_2$  with 0.75




Best  $X_2$  split: 0.6610, 0.482  
 B  
 B

There is no point in splitting this node further since it contains only data from a single class  $\rightarrow$  return it as a leaf node with output 'A'










Final decision tree  
 Given an input  $(X, Y) \rightarrow$   
 Follow the tree down to a  
 leaf.

Return corresponding  
 output class for this leaf

Example  $(X, Y) = (0.5, 0.5)$




# Basic Questions

- How to choose the attribute/value to split on at each level of the tree?
-  • When to stop splitting? When should a node be declared a leaf?
-  • If a leaf node is impure, how should the class label be assigned?
- If the tree is too large, how can it be pruned?

# When to stop splitting? Common strategies:

## 1. Pure and impure leave nodes

- All points belong to the same class; OR
- All points from one class completely overlap with points from another class (i.e., same attributes)
  - Output majority class as this leaf's label



## 2. Node contains points fewer than some threshold

## 3. Node purity is higher than some threshold

## 4. Further splits provide no improvement in *training loss*

$$(loss(T) \leq loss(T_L) + loss(T_R))$$



## Decision Tree Algorithm (Continuous Attributes)

- LearnTree( $X, Y$ )
  - Input:
    - Set  $X$  of  $R$  training vectors, each containing the values  $(x_1, \dots, x_M)$  of  $M$  attributes  $(X_1, \dots, X_M)$
    - A vector  $Y$  of  $R$  elements, where  $y_j$  = class of the  $j^{\text{th}}$  datapoint
  - If all the datapoints in  $X$  have the same class value  $y$ 
    - Return a leaf node that predicts  $y$  as output
  - If all the datapoints in  $X$  have the same attribute value  $(x_1, \dots, x_M)$ 
    - Return a leaf node that predicts the majority of the class values in  $Y$  as output
  - Try all the possible attributes  $X_j$  and threshold  $t$  and choose the one,  $j^*$ , for which  $\text{IG}(Y|X_j, t)$  is maximum
  - $X_L, Y_L$  = set of datapoints for which  $x_{j^*} < t$  and corresponding classes
  - $X_H, Y_H$  = set of datapoints for which  $x_{j^*} \geq t$  and corresponding classes
  - Left Child  $\leftarrow \text{LearnTree}(X_L, Y_L)$
  - Right Child  $\leftarrow \text{LearnTree}(X_H, Y_H)$

## Decision Tree Algorithm (Discrete Attributes)


- LearnTree( $X, Y$ )
  - Input:
    - Set  $X$  of  $R$  training vectors, each containing the values  $(x_1, \dots, x_M)$  of  $M$  attributes  $(X_1, \dots, X_M)$
    - A vector  $Y$  of  $R$  elements, where  $y_j$  = class of the  $j^{\text{th}}$  datapoint
  - If all the datapoints in  $X$  have the same class value  $y$ 
    - Return a leaf node that predicts  $y$  as output
  - If all the datapoints in  $X$  have the same attribute value  $(x_1, \dots, x_M)$ 
    - Return a leaf node that predicts the majority of the class values in  $Y$  as output
  - Try all the possible attributes  $X_j$  and choose the one,  $j^*$ , for which  $IG(Y|X_j)$  is maximum
  - For every possible value  $v$  of  $X_{j^*}$ :
    - $X_v, Y_v$  = set of datapoints for which  $x_{j^*} = v$  and corresponding classes
    - $\text{Child}_v \leftarrow \text{LearnTree}(X_v, Y_v)$

# Decision Trees So Far

- Given  $N$  observations from training data, each with  $D$  attributes  $X$  and a class attribute  $Y$ , construct a sequence of tests (decision tree) to predict the class attribute  $Y$  from the attributes  $X$
- Basic strategy for defining the tests (“when to split”) → maximize the information gain on the training data set at each node of the tree
- Problems (next):
  - Computational issues → How expensive is it to compute the IG
  - The tree will end up being much too big → *pruning*
  - Evaluating the tree on training data is dangerous → *overfitting*

# Basic Questions

- How to choose the attribute/value to split on at each level of the tree?
- When to stop splitting? When should a node be declared a leaf?
- If a leaf node is impure, how should the class label be assigned?
- If the tree is too large, how can it be pruned?



# What will happen if a tree is too large?

Overfitting

High variance

Instability in predicting test data

# How to avoid overfitting?

- Acquire more training data
- Remove irrelevant attributes (manual process – not always possible)
- Grow full tree, then post-prune
- Ensemble learning

# Reduced-Error Pruning

Split data into training and validation sets

Grow tree based on training set

Do until further pruning is harmful:

1. Evaluate impact on validation set of pruning each possible node (plus those below it)
2. Greedily remove the node that most improves validation set accuracy



# How to decide to remove it a node using pruning

- Pruning of the decision tree is done by replacing a whole subtree by a leaf node.
- The replacement takes place if a decision rule establishes that the expected error rate in the subtree is greater than in the single leaf.

