

# James Webb Space Telescope





# The landscape of “mountains” and “valleys”





# A visual grouping of five galaxies



# CONVOLUTIONAL NEURAL NETWORK

Mahdi Roozbahani

Georgia Tech


**Great visualization tool:**

**<https://poloclub.github.io/cnn-explainer/>**

Slides are based on Ming Li (University of Waterloo – Deep learning part) with some modifications



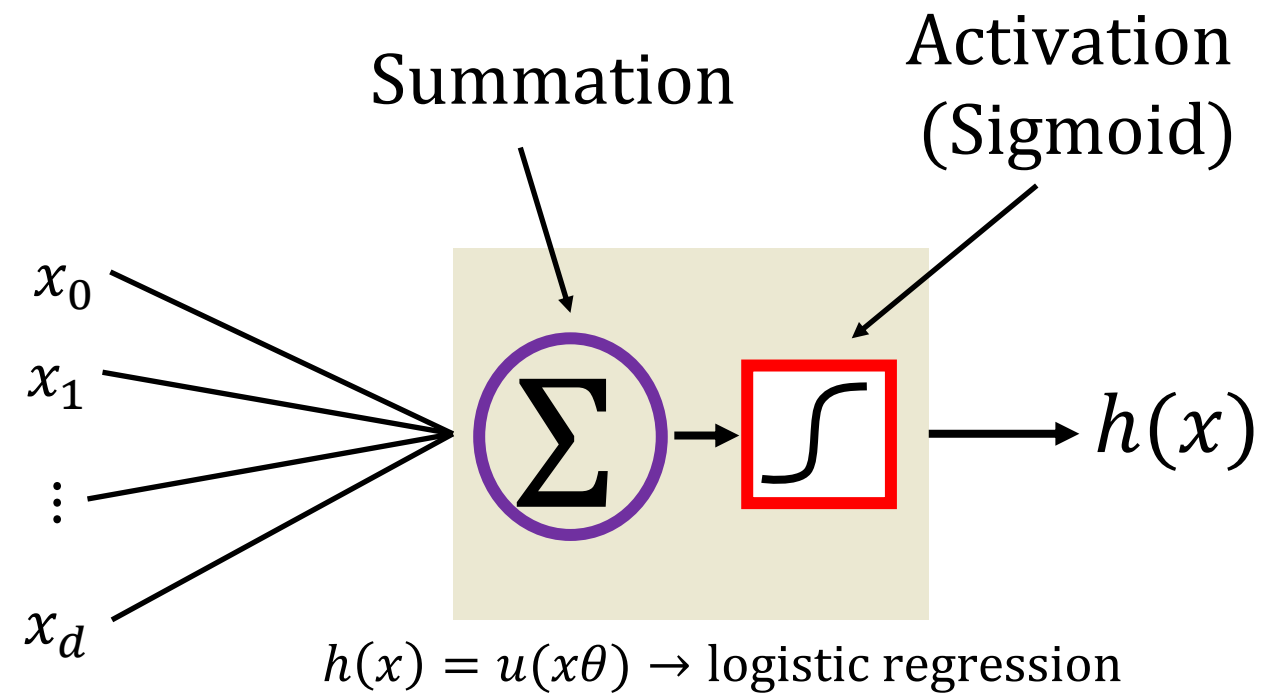
# Inspiration from Biological Neurons



The first drawing of a brain cells by Santiago Ramón y Cajal in 1899

**Neurons:** core components of brain and the nervous system consisting of



1. Dendrites that collect information from other neurons
2. An axon that generates outgoing spikes



$$\text{output} = \text{activation}(x\theta + b)$$

Name of the neuron	Activation function: $\text{activation}(z)$
Linear unit	$x\theta$
Threshold/sign unit	$\text{sign}(x\theta)$
Sigmoid unit	$\frac{1}{1 + \exp(x\theta)}$
Rectified linear unit (ReLU)	$\max(0, x\theta)$
Tanh unit	$\tanh(x\theta)$

# NN Regression




# FACIAL RECOGNITION

Deep-learning neural networks use layers of increasingly complex rules to categorize complicated shapes such as faces.



Layer 1: The computer identifies pixels of light and dark.



Layer 2: The computer learns to identify edges and simple shapes.



Layer 3: The computer learns to identify more complex shapes and objects.




Layer 4: The computer learns which shapes and objects can be used to define a human face.



# Smaller Network: CNN


- We know it is good to learn a small model.
- From this fully connected model, do we really need all the edges?
- Can some of these be shared?



# Consider learning an image:

- Some patterns are much smaller than the whole image




Can represent a small region with fewer parameters






Same pattern appears in different places:  
They can be compressed!

What about training a lot of such “small” detectors  
and each detector must “move around”.



# A convolutional layer

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that does convolutional operation.





# Convolution

**These are the network parameters to be learned.**

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮ ⋮

Each filter detects a small pattern (3 x 3).

# Convolution


stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1





# Convolution


If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

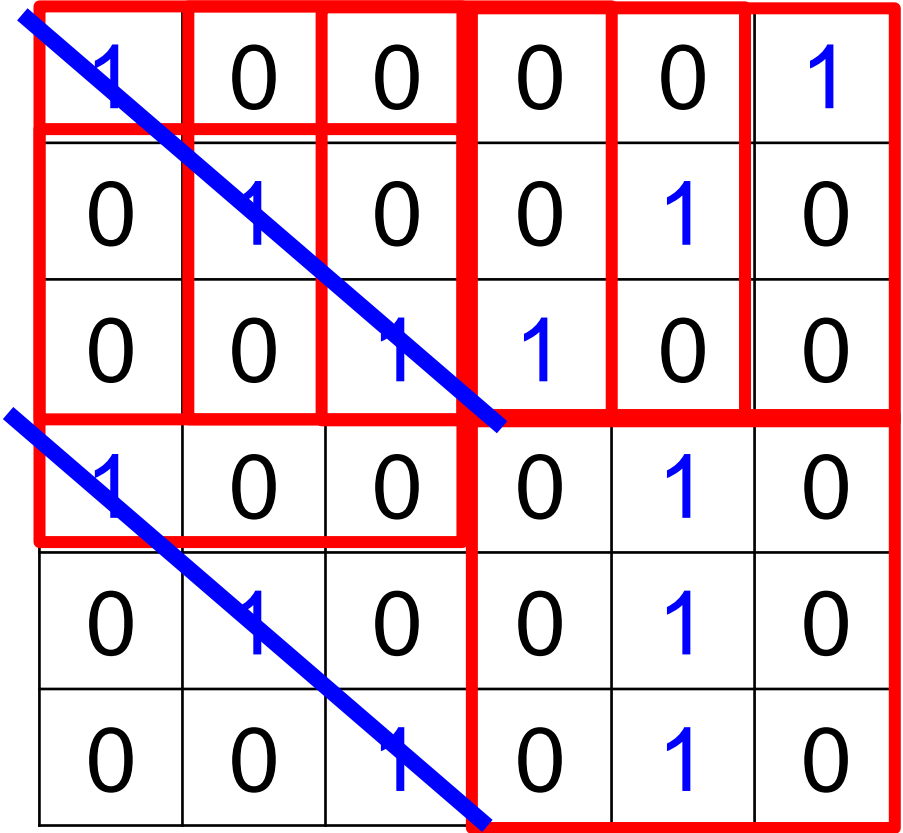
1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



# Convolution

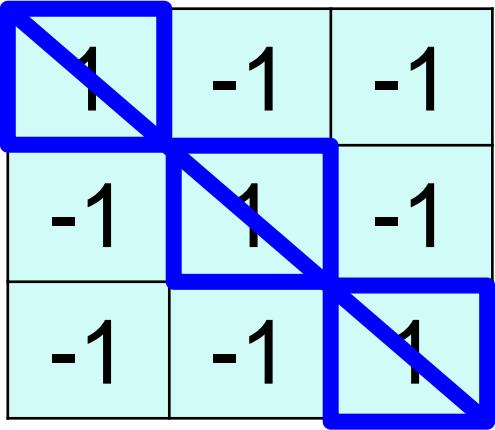
stride=1



A 6x6 grid of numbers representing an image. The values are: Row 1: 1, 0, 0, 0, 0, 1; Row 2: 0, 1, 0, 0, 1, 0; Row 3: 0, 0, 1, 1, 0, 0; Row 4: 1, 0, 0, 0, 1, 0; Row 5: 0, 1, 0, 0, 1, 0; Row 6: 0, 0, 1, 0, 1, 0. A 3x3 red bounding box highlights the top-left corner (rows 1-3, columns 1-3). A blue diagonal line runs from the top-left cell (1,1) to the bottom-right cell (6,6).

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

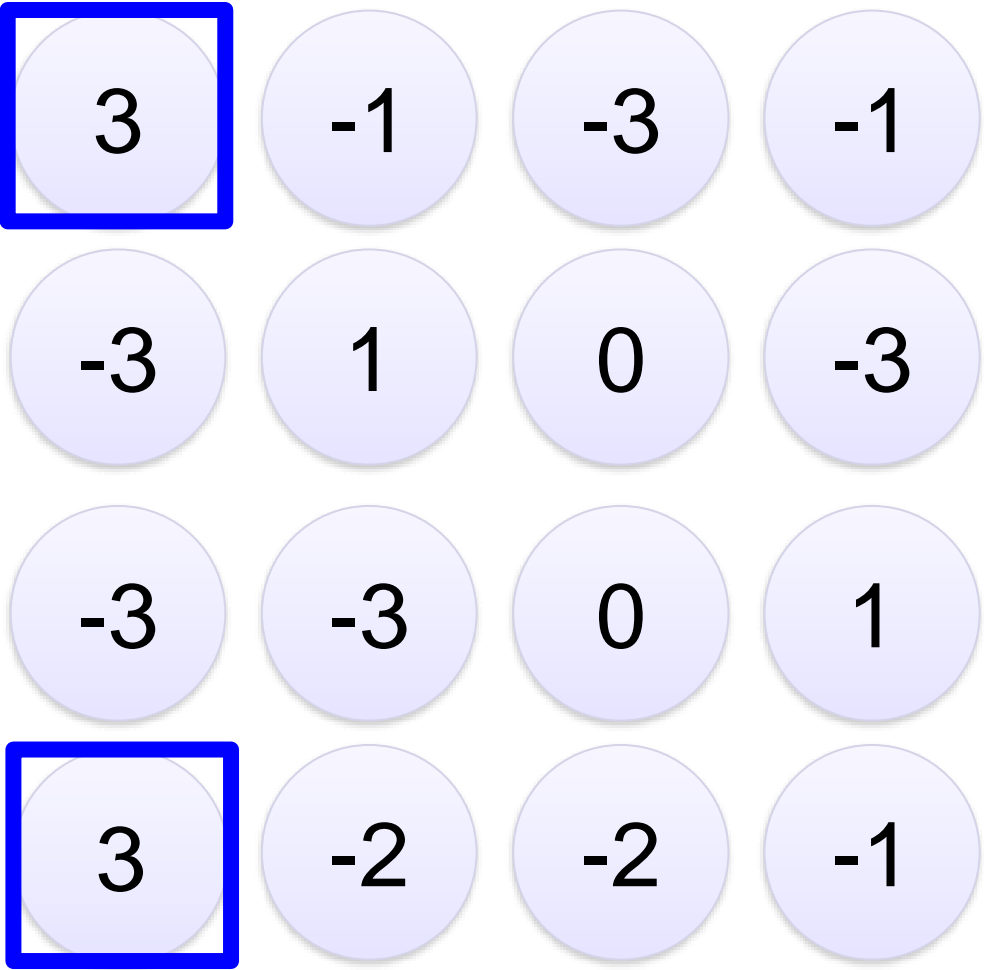
6 x 6 image



A 3x3 grid of numbers representing a filter. The values are: Row 1: 1, -1, -1; Row 2: -1, 1, -1; Row 3: -1, -1, 1. A blue border surrounds the entire grid, and a blue diagonal line runs from the top-left cell (1,1) to the bottom-right cell (3,3).

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



A 4x4 grid of circles containing numbers. The values are: Row 1: 3, -1, -3, -1; Row 2: -3, 1, 0, -3; Row 3: -3, -3, 0, 1; Row 4: 3, -2, -2, -1. The circles in the first column (top-left and bottom-left) are highlighted with blue square borders.

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1



# Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

Repeat this for each filter



Two 4 x 4 images  
Forming 2 x 4 x 4 matrix

# Color image: RGB 3 channels





# Convolution v.s. Fully Connected

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

image

1	-1	-1
-1	1	-1
-1	-1	1

-1	1	-1
-1	1	-1
-1	1	-1




convolution

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3

Fully-  
connected

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0




Conventional  
Fully Connected  
layers  
(FC layers)

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


6 x 6 image

1	1	
2	0	
3	0	
4	0	
5	0	
6	1	
7	0	
8	1	
⋮		⋮
31	0	
32	0	
33	1	
34	0	
35	1	
36	0	


features    1<sup>st</sup> hidden layer



fewer parameters!








6 x 6 image

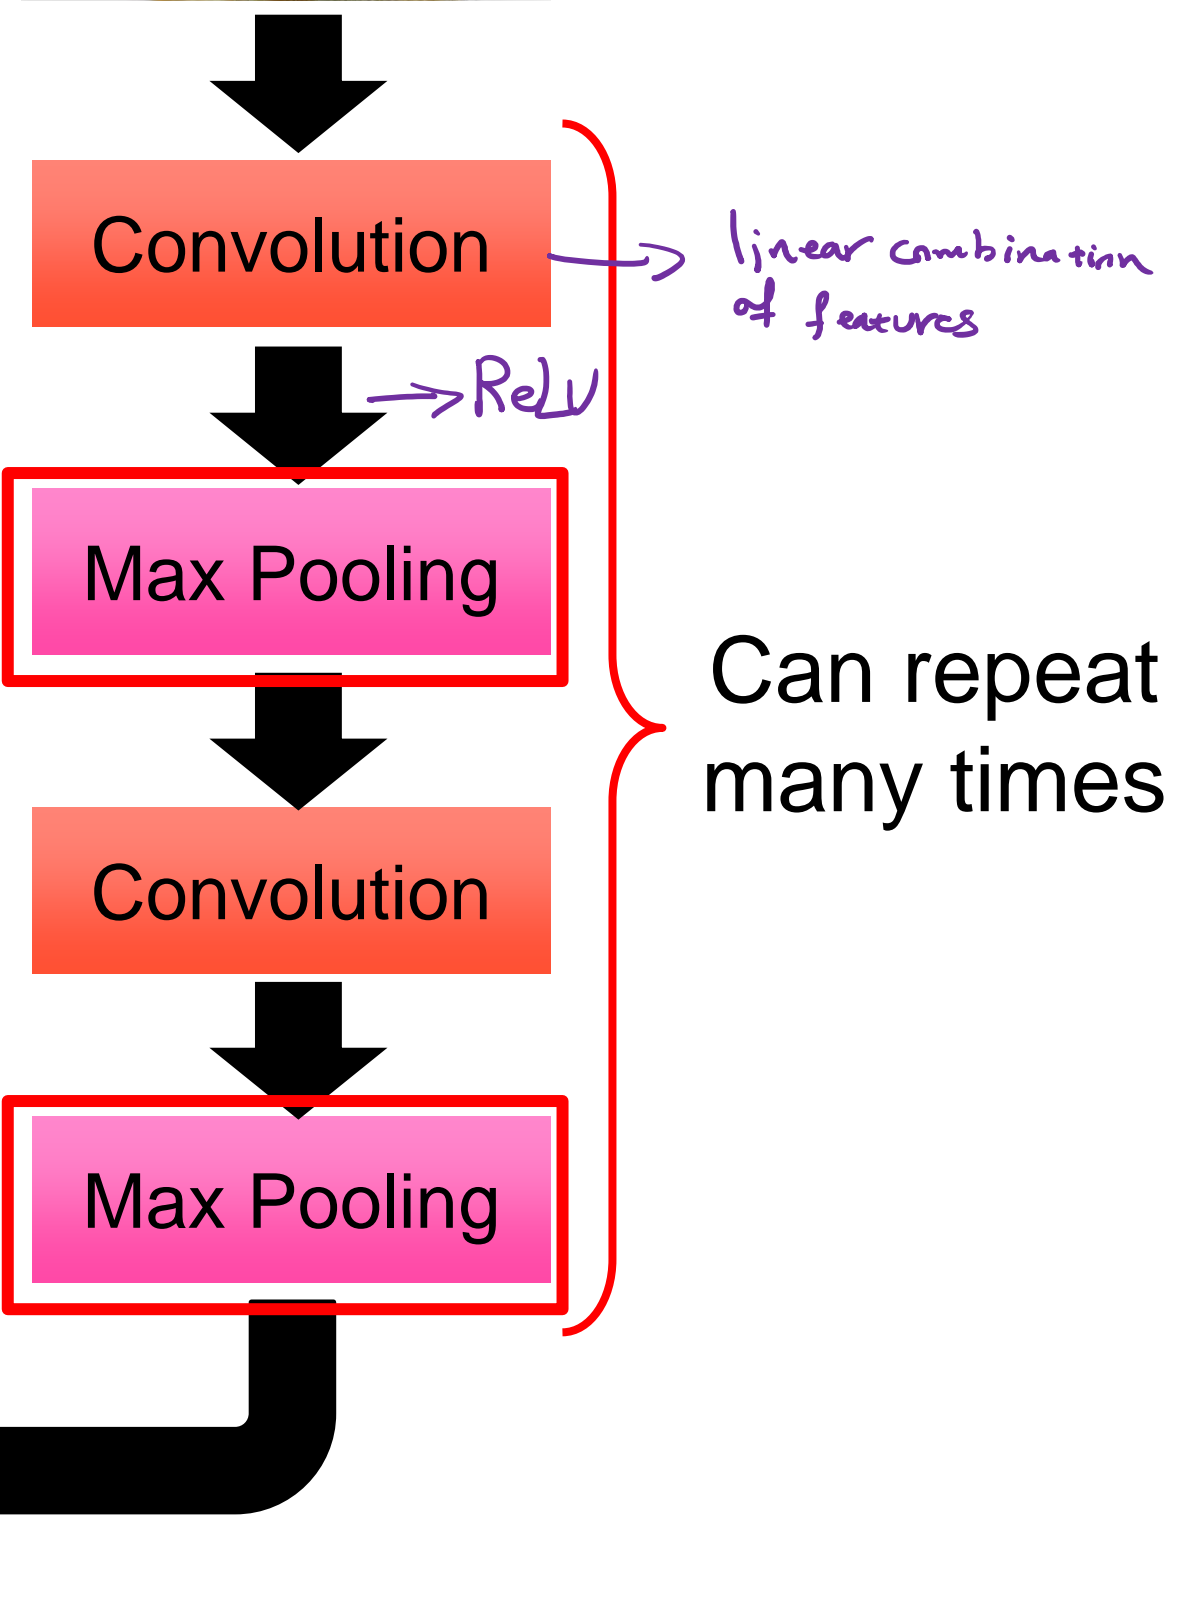

Fewer parameters

Even fewer parameters



# The whole CNN

cat dog .....



# Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3



# Why Pooling

- Subsampling pixels will not change the object  
bird




Subsampling



bird

We can subsample the pixels to make image smaller



fewer parameters to characterize the image


# A CNN compresses a fully connected network in three ways:

- Reducing number of connections
- Shared weights on the edges
- Max pooling further reduces the complexity

# Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


6 x 6 image




Conv



Max  
Pooling




New image  
but smaller



2 x 2 image


Each filter  
is a channel

# The whole CNN



Smaller than the original image

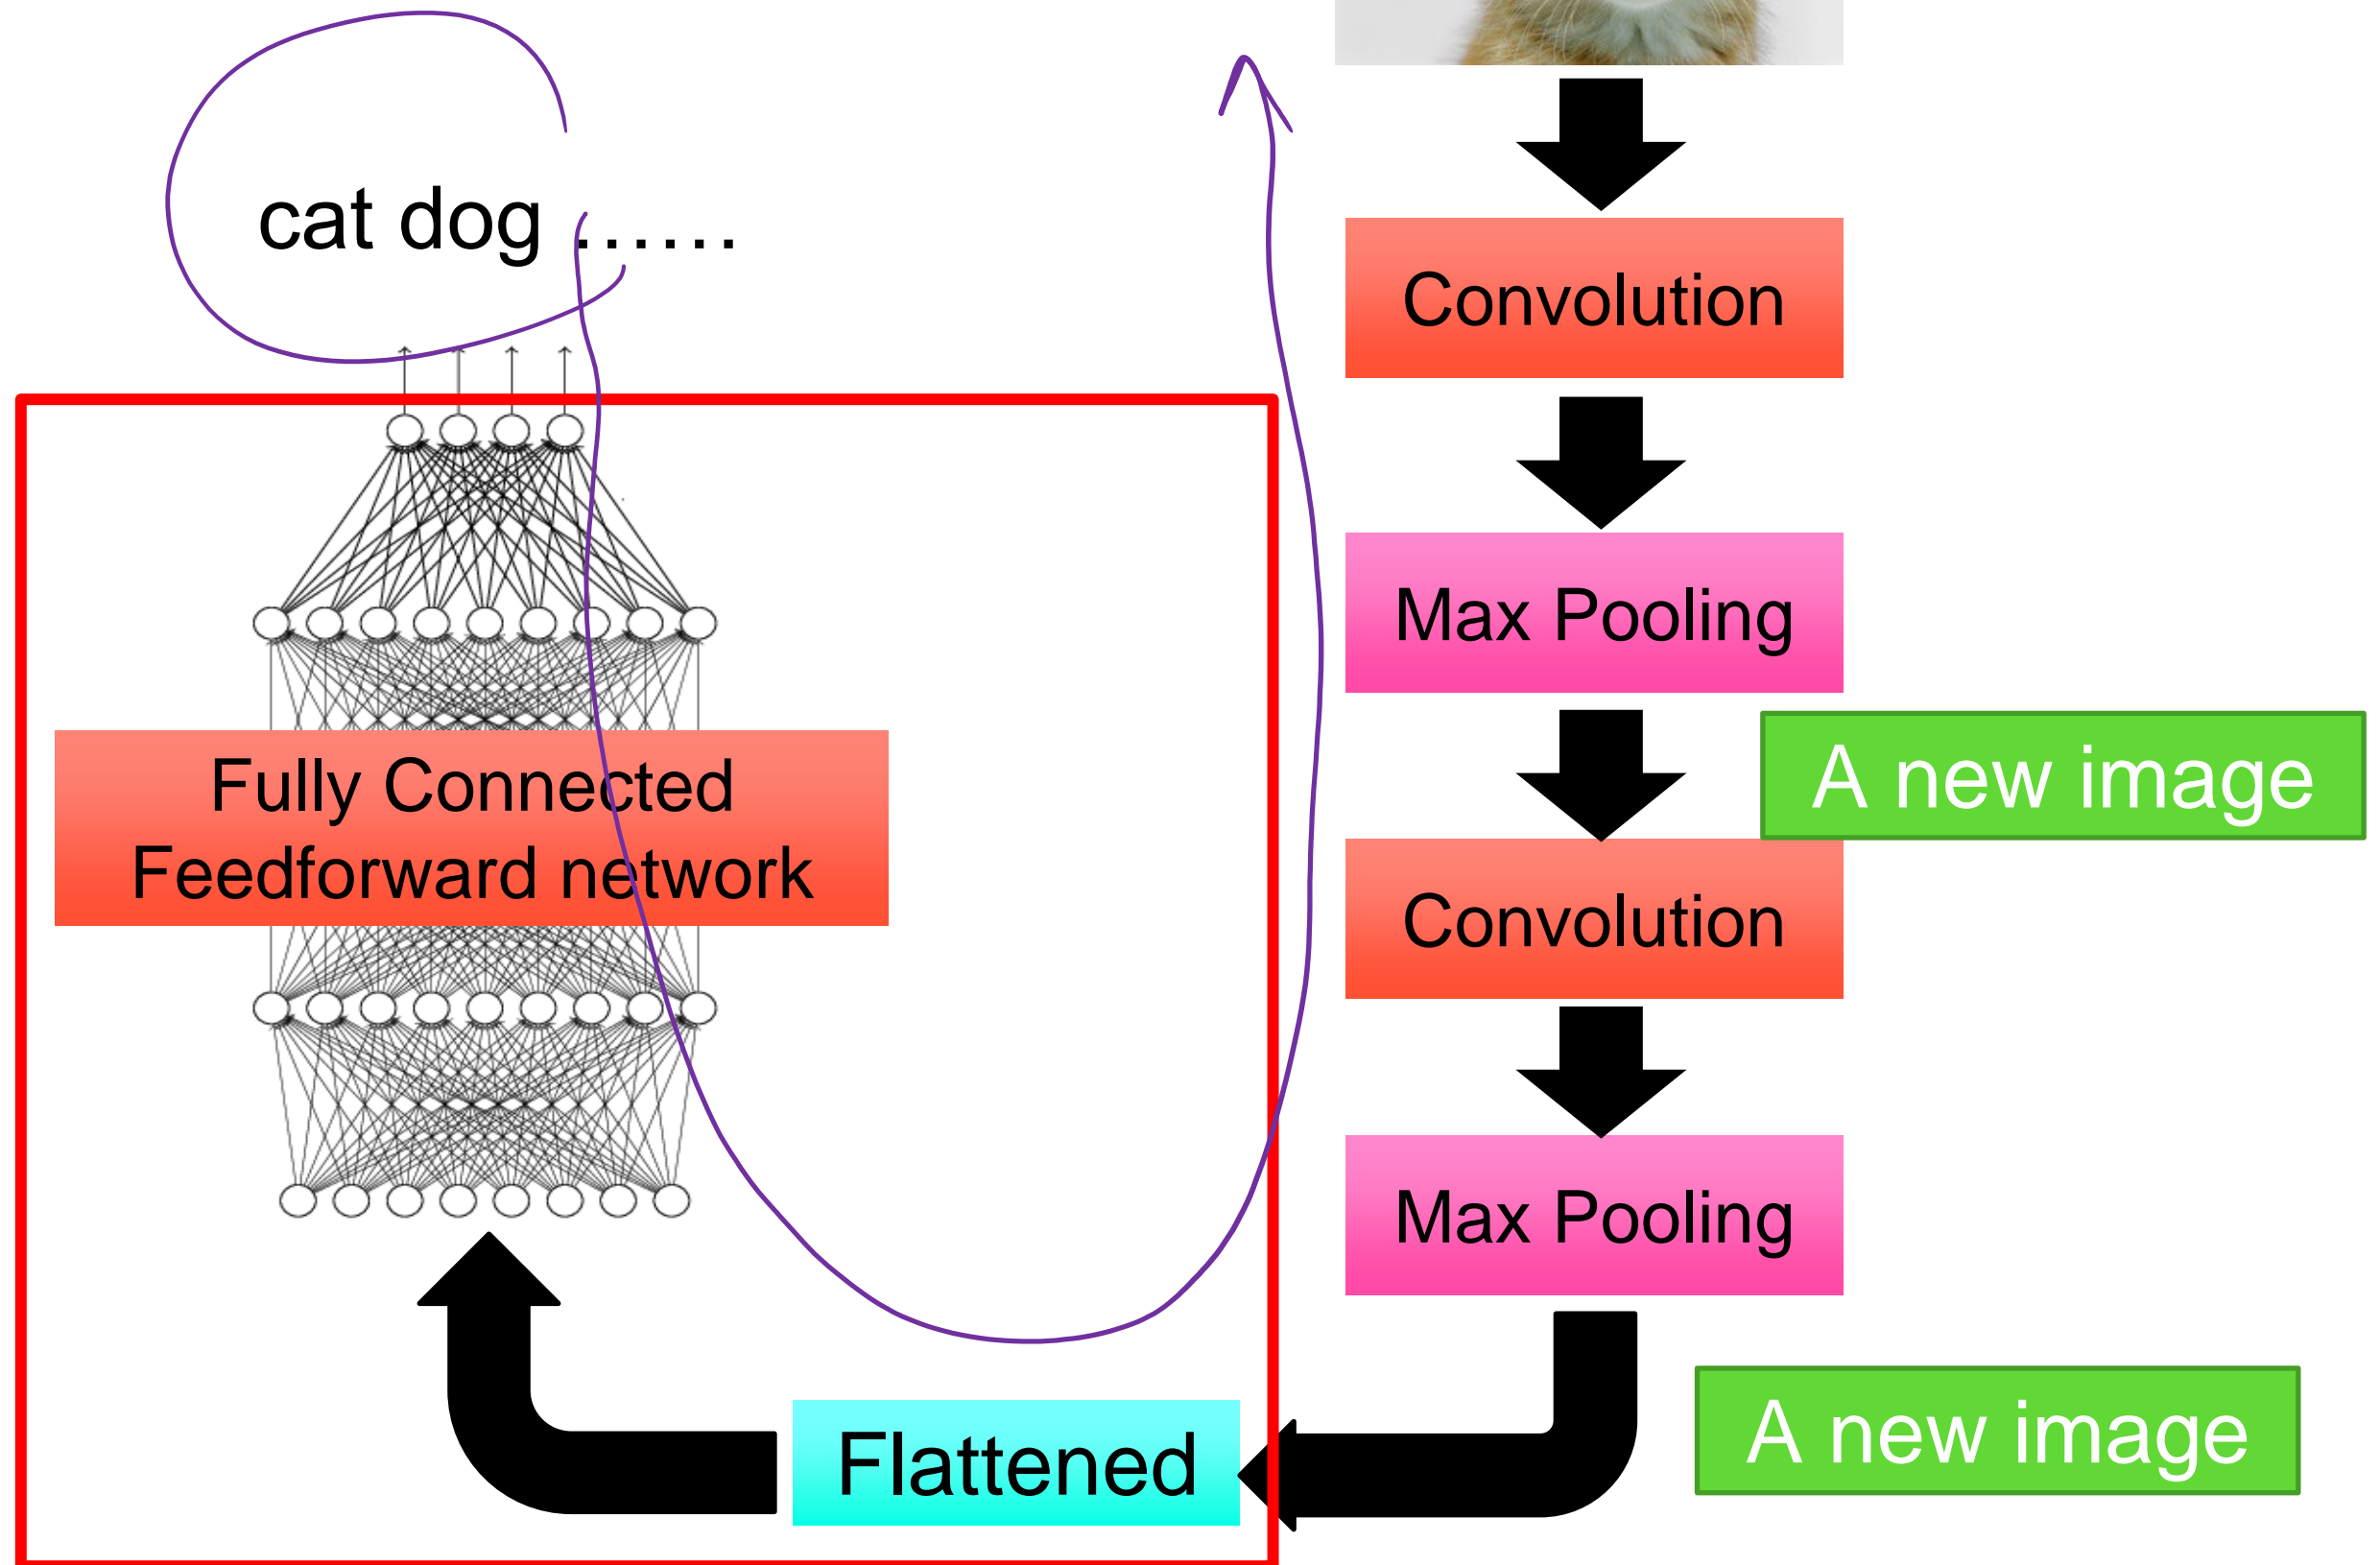
The number of channels is the number of filters




Can repeat many times



# The whole CNN



# Flattening



# CNN in Keras

Only modified the *network structure* and *input format (vector -> 3-D tensor)*

```
model2.add( Convolution2D( 25,3,3,  
                           input_shape=(28,28,1)) )
```

1	-1	-1	1	-1
-1	1	-1	1	-1
-1	-1	-1	1	-1


There are  
**25 3x3**  
filters.

Input\_shape = ( 28 , 28 , 1)

28 x 28 pixels

1: black/white, 3: RGB

```
model2.add(MaxPooling2D((2,2)))
```



input

Convolution

-> ReLU

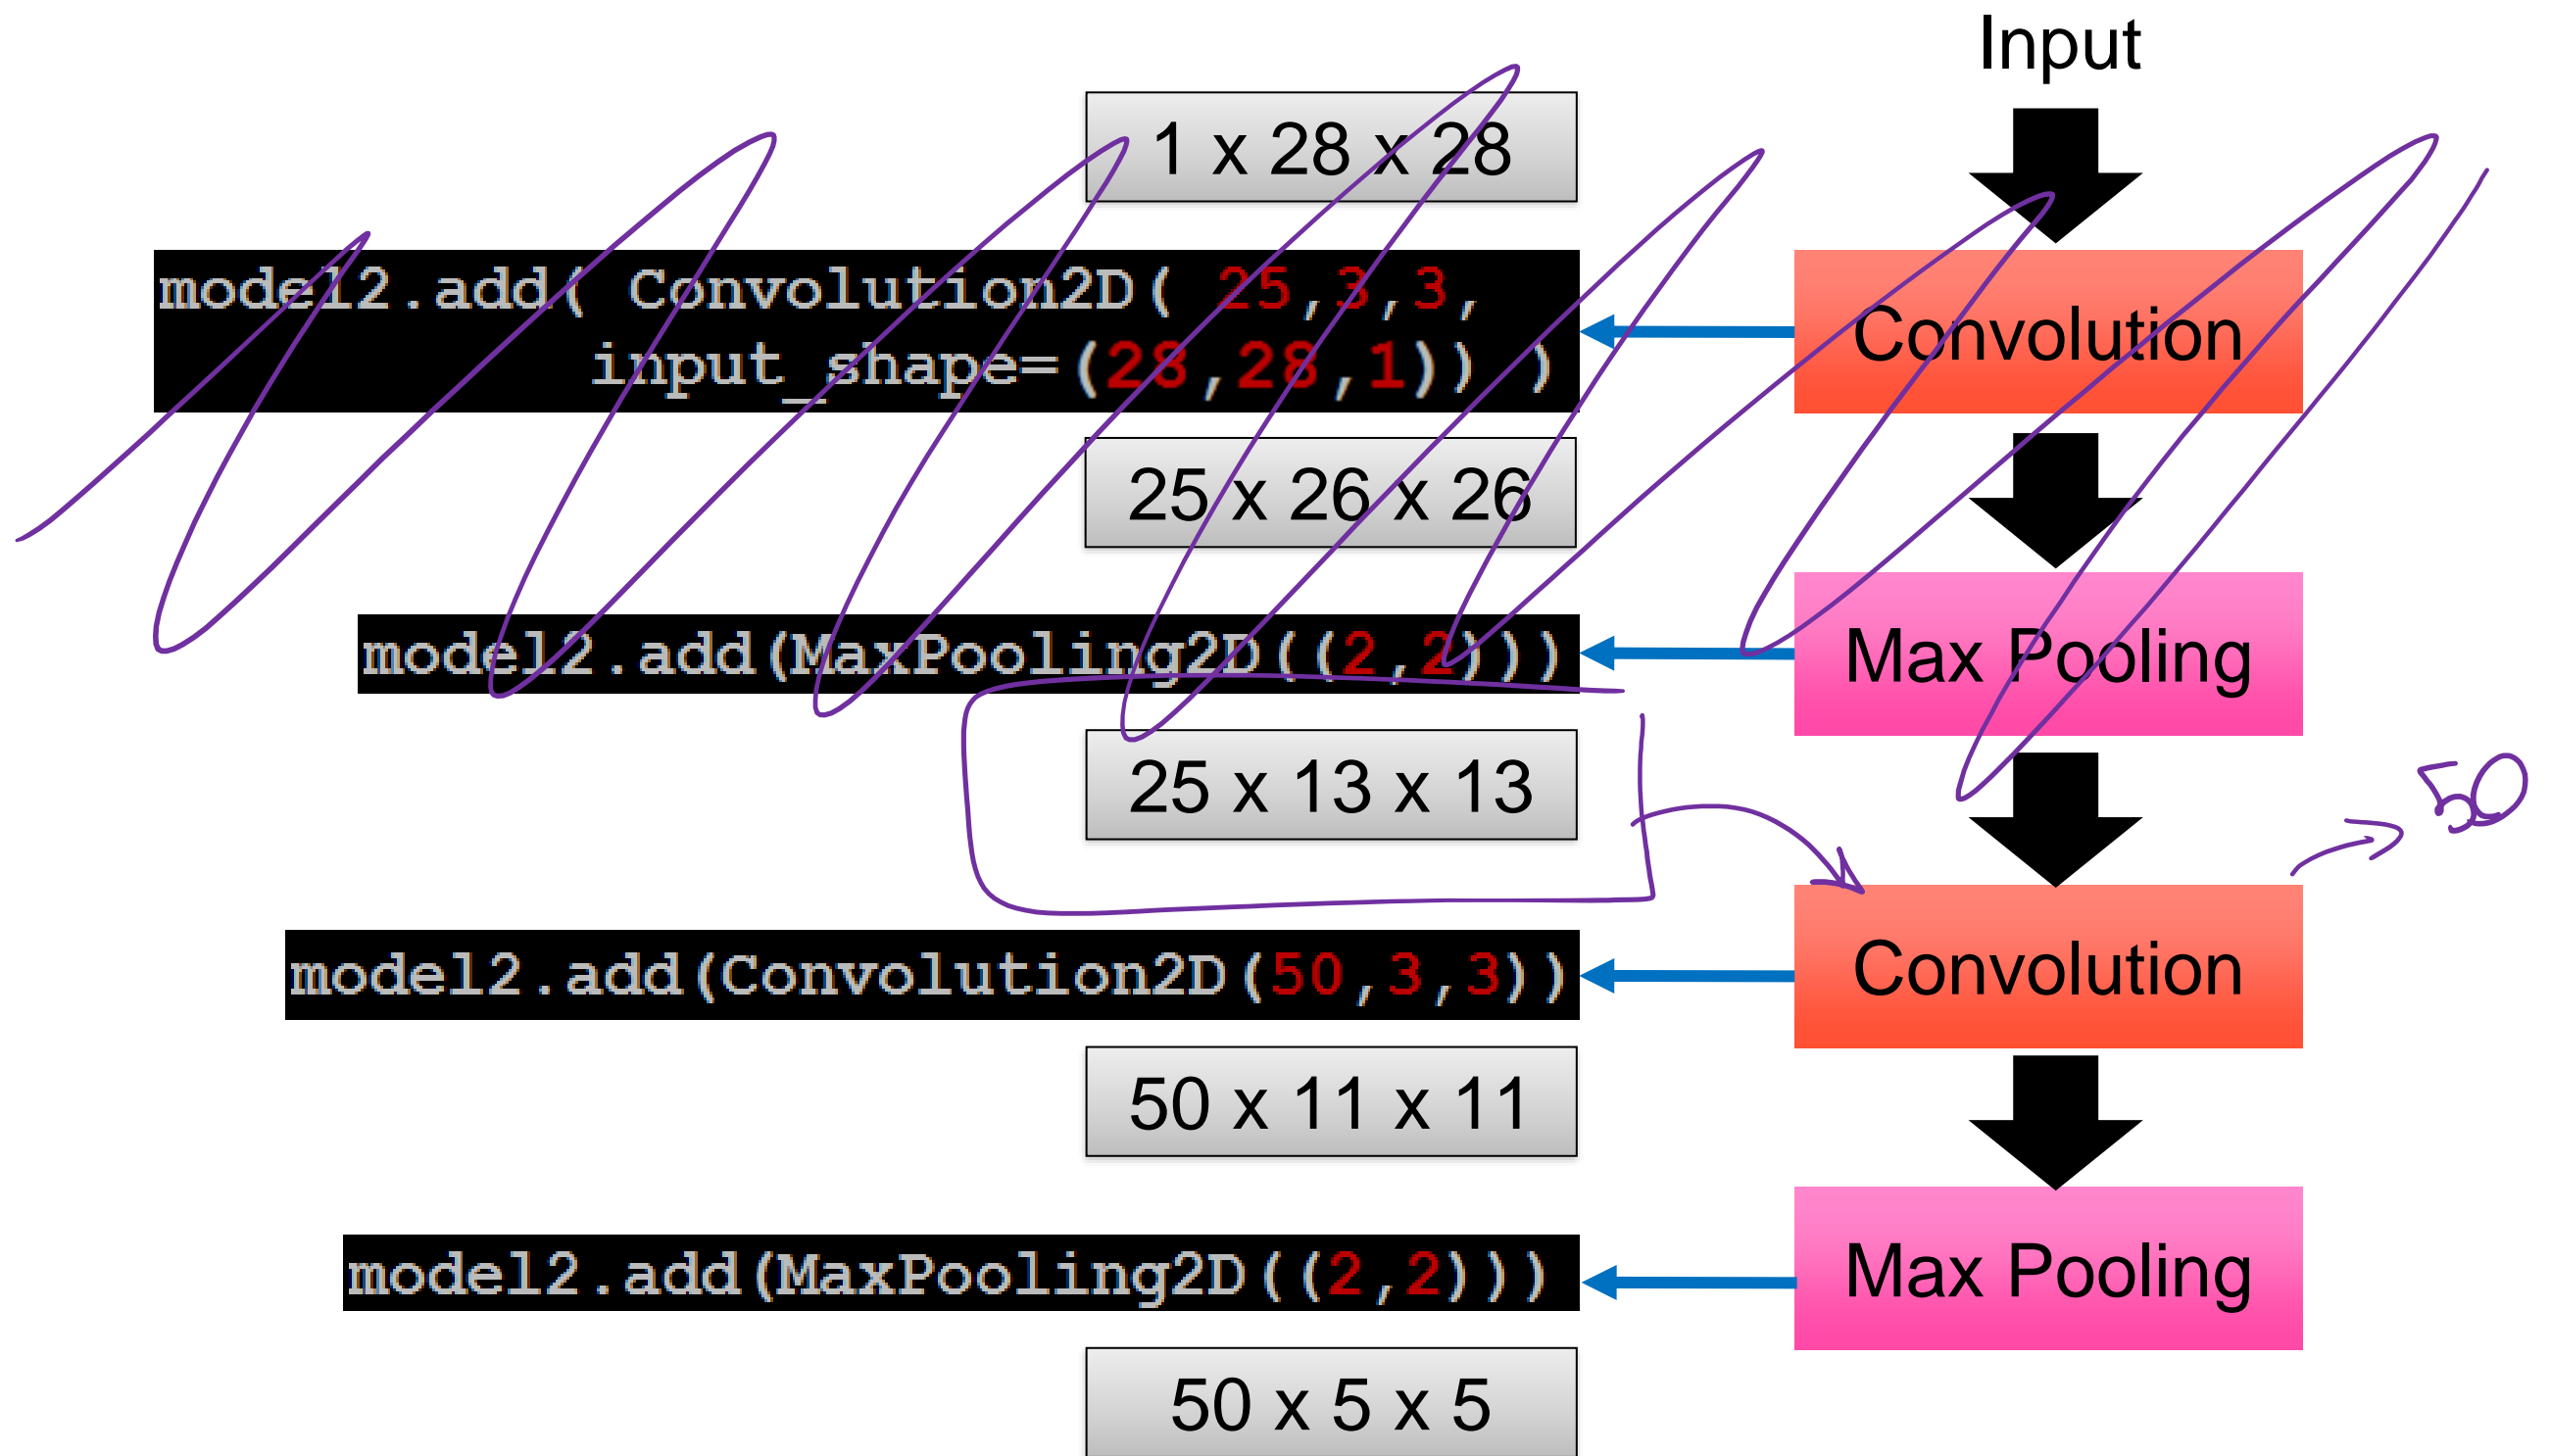
Max Pooling

Convolution

Max Pooling

# CNN in Keras

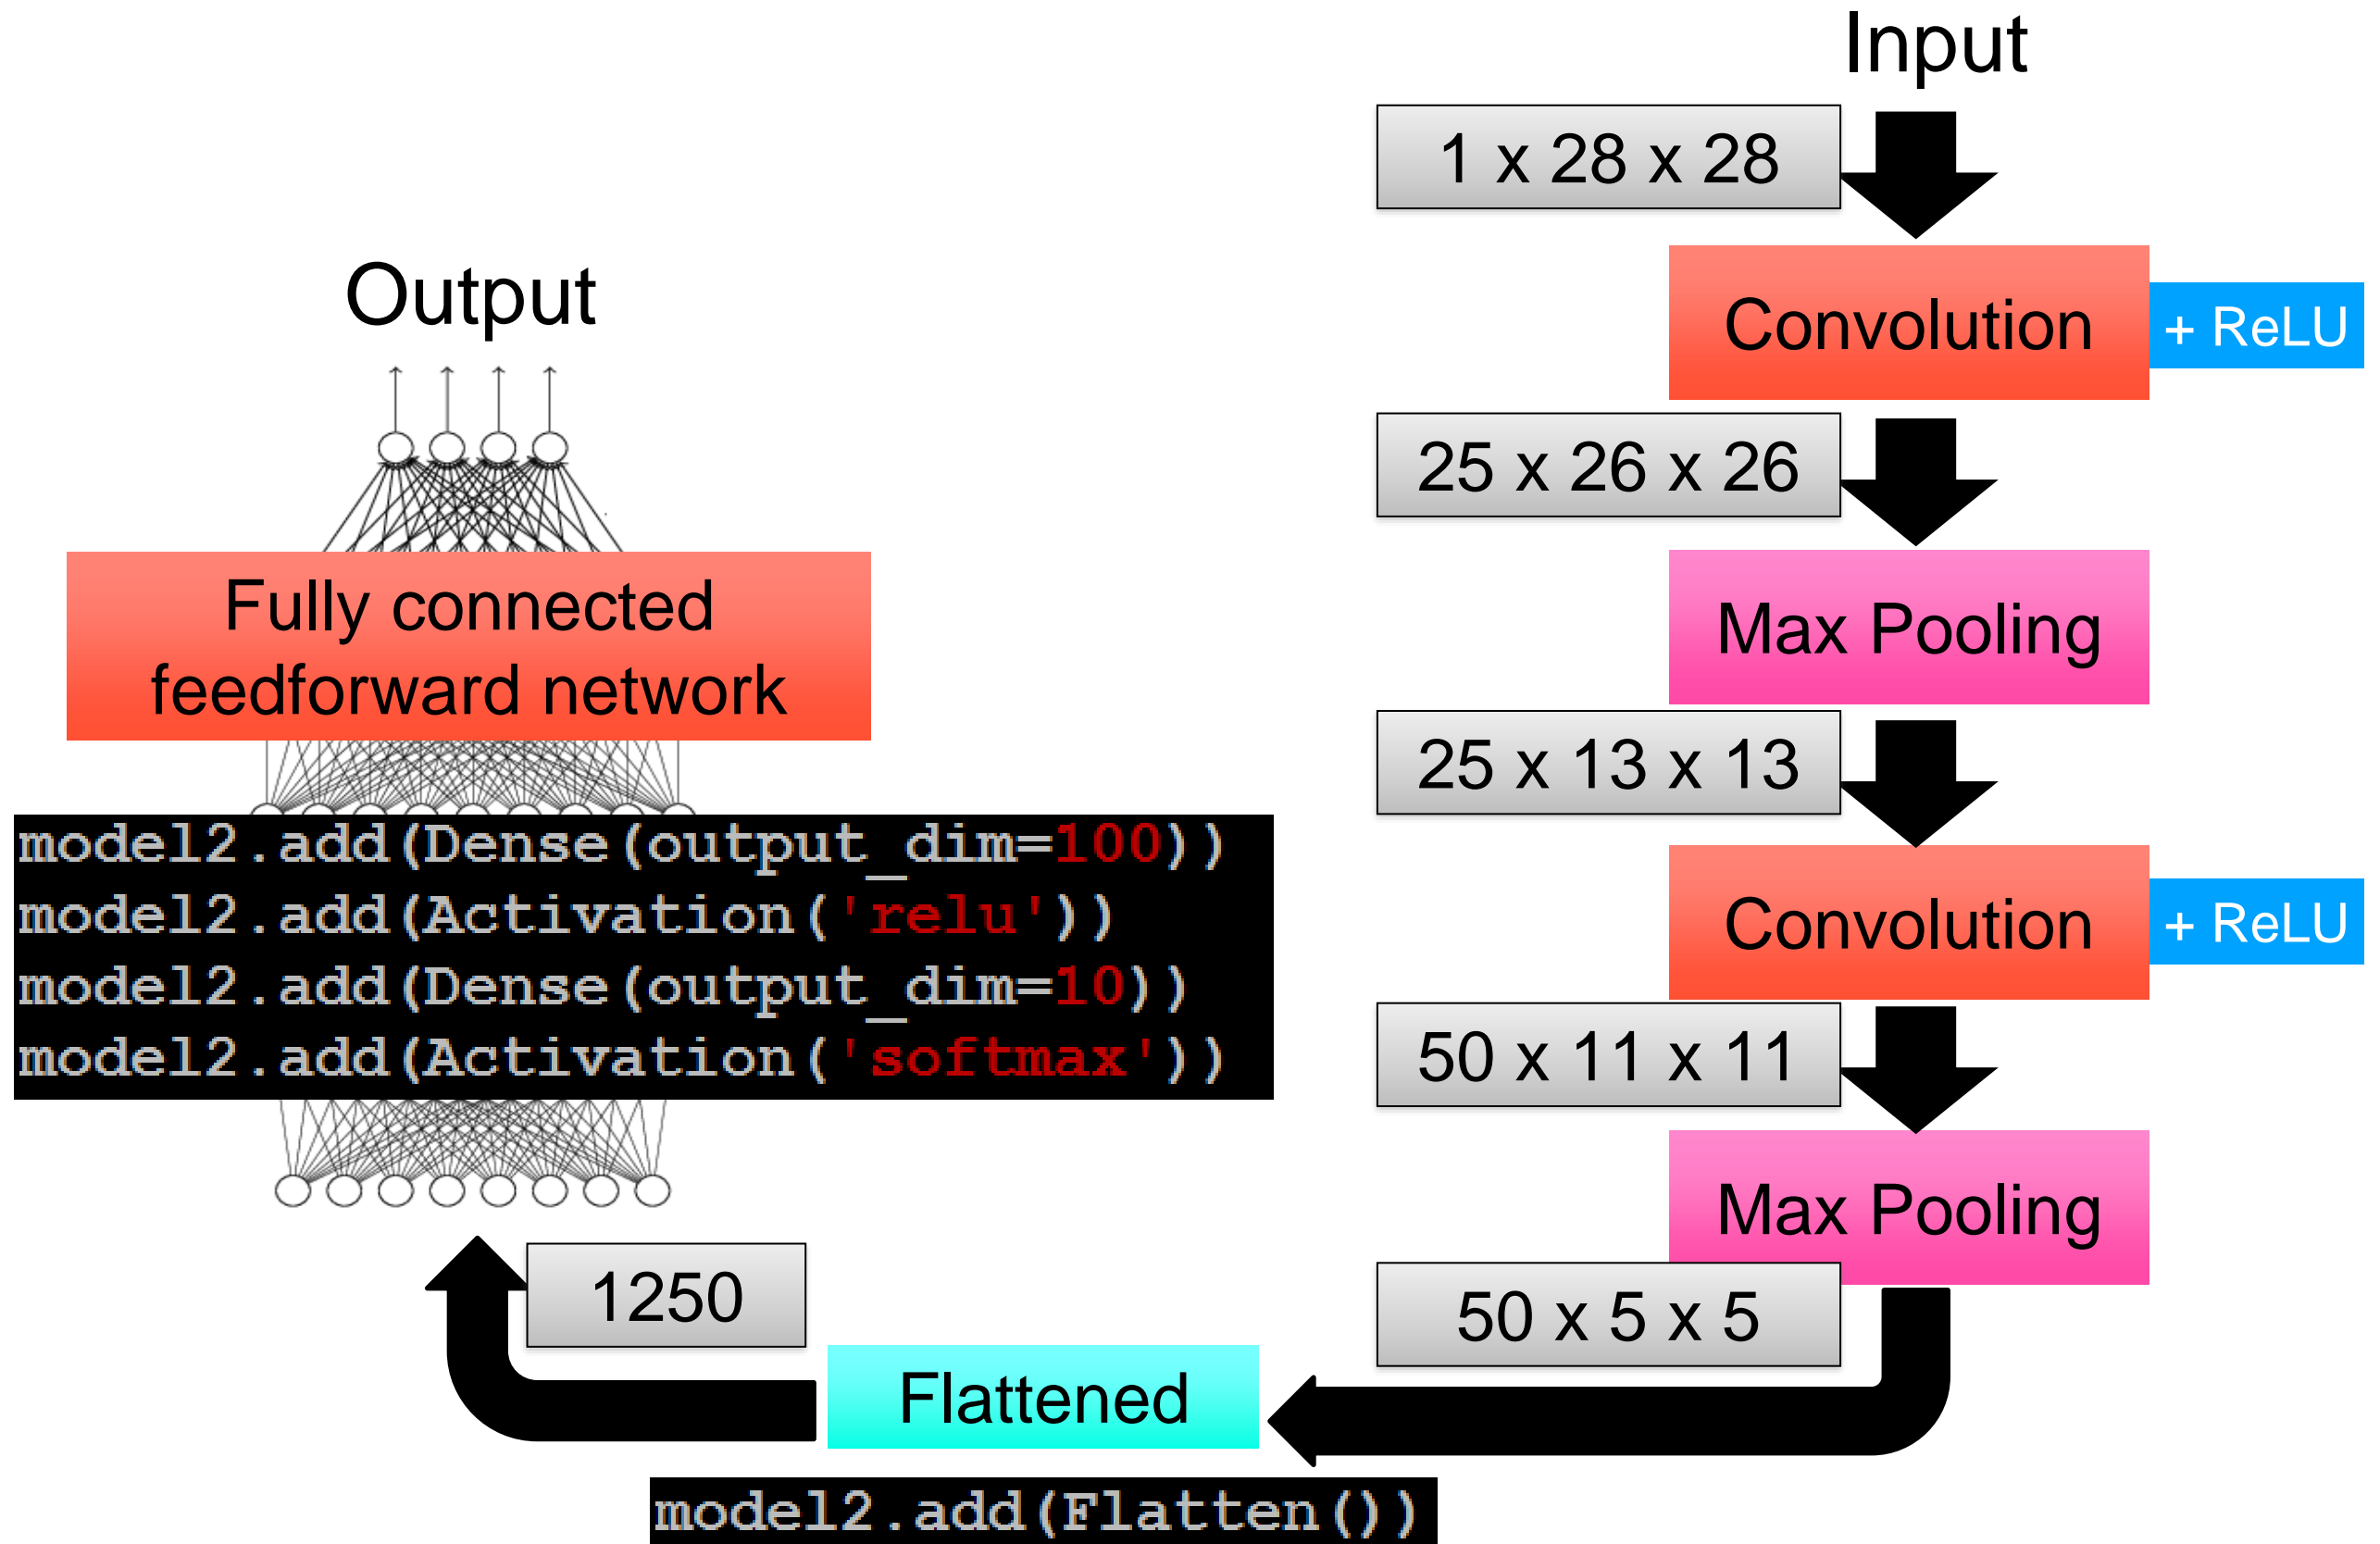
Only modified the *network structure* and *input format* (vector -> 3-D array)





# CNN in Keras

Only modified the *network structure* and *input format* (vector -> 3-D array)




# Number of Parameters

$25 \times 3 \times 3 + 25$  parameters


25 filters - Conv1

25:  $3 \times 3$


Gray scaled-image




Convolved result for 25 filters



Result after Max Pooling





Result after Max Pooling



50:  $3 \times 3 \times 25$

50 filters - Conv2



Convolved result for 50 filters



$50 \times 3 \times 3 \times 25 + 50$  parameters

# neurons after flattening:  $50 \times 5 \times 5 = 1250$

Flattening



10 CNN Architecture